

mcvis: multicollinearity visualisation

https://kevinwang09.github.io/pres/mcvis_talk


Kevin Y. X. Wang

5 December 2019, Adelaide



Acknowledgement

This is joint work with Chen Lin (Fudan University) and Prof Samuel Mueller (University of Sydney).



Cricketers' career batting statistics

- Cricket is a bat-and-ball game.
- The aim of a batsman is to score as many **runs** as possible before getting **out**.


```
glimpse(X)
```

```
## Observations: 810
## Variables: 8
## $ log_runs    <dbl> 2.20, 1.56, 2.84, 2.68, 2.01, 3.21, 2.03, 2.65, 3.13, 2.68,...
## $ log_outs    <dbl> 1.040, 0.778, 1.410, 1.320, 1.230, 1.610, 1.000, 1.490, 1.7...
## $ log_ave     <dbl> 1.160, 0.778, 1.430, 1.360, 0.778, 1.600, 1.030, 1.160, 1.4...
## $ log_fours   <dbl> 1.280, 0.301, 1.830, 1.830, 1.040, 2.160, 1.110, 1.650, 2.0...
## $ log_sixes   <dbl> 0.000, 0.000, 0.477, 0.845, 0.301, 0.602, 0.000, 0.477, 0.6...
## $ log_ducks   <dbl> 0.699, 0.477, 0.602, 0.602, 1.040, 0.778, 0.602, 0.845, 0.6...
## $ log_hs      <dbl> 2.07, 1.26, 2.02, 2.00, 1.41, 2.10, 1.48, 1.52, 2.10, 1.95,...
## $ log_100     <dbl> 0.301, 0.000, 0.301, 0.301, 0.000, 0.699, 0.000, 0.000, 0.4...
```

Interesting feature in this data

There is a causal relationship:

$$\text{batting ave} = \frac{\text{runs}}{\text{no. of outs}}, \quad \text{or equivalently,} \quad \log_{\text{runs}} = \log_{\text{ave}} + \log_{\text{outs}}.$$



What is multi-collinearity (MC)?

MC occurs when columns of X are linear dependent (exactly or approximately).

```
M1 = lm(log_100 ~ ., data = X)
broom::tidy(M1)
```


```
## # A tibble: 8 x 5
##   term          estimate std.error statistic  p.value
##   <chr>          <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept) -0.365      0.0902    -4.05  5.67e- 5
## 2 log_runs    -1.92       1.95     -0.984  3.25e- 1
## 3 log_outs     1.61       1.96      0.826  4.09e- 1
## 4 log_ave      1.84       1.96      0.943  3.46e- 1
## 5 log_fours     0.647     0.0969     6.68  4.58e-11
## 6 log_sixes     0.131     0.0264     4.96  8.57e- 7
## 7 log_ducks     0.00357    0.0497     0.0718 9.43e- 1
## 8 log_hs      -0.0187    0.0753    -0.248  8.04e- 1
```

Consequence of multi-collinearity

- We will proceed with rounding all variables to 3 significant figures.

	Include all			Remove log_runs			Remove log_ave		
<i>Predictors</i>	<i>Estimates</i>	<i>std. Error</i>	<i>p</i>	<i>Estimates</i>	<i>std. Error</i>	<i>p</i>	<i>Estimates</i>	<i>std. Error</i>	<i>p</i>
(Intercept)	-0.37	0.09	<0.001	-0.37	0.09	<0.001	-0.36	0.09	<0.001
log_runs	-1.92	1.95	0.325				-0.08	0.12	0.491
log_outs	1.61	1.96	0.409	-0.31	0.11	0.004	-0.23	0.10	0.019
log_ave	1.84	1.96	0.346	-0.08	0.12	0.530			
log_fours	0.65	0.10	<0.001	0.64	0.10	<0.001	0.65	0.10	<0.001
log_sixes	0.13	0.03	<0.001	0.13	0.03	<0.001	0.13	0.03	<0.001
log_ducks	0.00	0.05	0.943	0.00	0.05	0.922	0.00	0.05	0.934
log_hs	-0.02	0.08	0.804	-0.02	0.08	0.811	-0.02	0.08	0.837

High correlation \neq multicollinearity



- By definition, it is the linear combination of variables that causes MC.
- The causal variables are not the most highly correlated.
- Thus, identifying high correlation does not always identify sources of MC.

Diagnosis of multicollinearity requires specialised statistics.

Existing methods

1. Variance inflation factors (VIFs)

Introduced in Marquardt (1970) and elsewhere:

$$VIF_j = \frac{1}{1 - R_j^2}, \quad j = 1, \dots, p,$$

where R_j^2 is the coefficient of determination when the x_j independent variable is treated as a response variable against the remaining $p - 1$ independent variables.

A **larger** value of VIF_j implies x_j can be highly predicted by other variables, and thus implies higher cause of MC by that variable.

```
M1 = lm(log_100 ~ ., data = X)
M1 %>% car::vif() %>% round(2)
```

```
##  log_runs  log_outs  log_ave log_fours log_sixes log_ducks  log_hs
##  23995.96  11410.15  4666.15   55.60     2.53     3.99   12.17
```

- Using a threshold of 5 as suggested by Sheather (2009), 5 MC-causing variables are identified.

2. Eigenvalues of $X^\top X$

Eigenvalues of the "uncentered covariance matrix" $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_p \geq 0$ offers a more linear algebra interpretation of MC.

A **smaller** value of λ_p produces a matrix determinant closer to 0, which implies linear dependence in X and thus MC (Stewart 1987).

```
Xmat = X %>% as.data.frame() %>% as.matrix() %>% scale()  
eigen = svd(t(Xmat) %*% Xmat)  
round(eigen$d, 3)
```

```
## [1] 4839.921  928.325  303.818  252.626   91.953   45.354    9.982    0.020
```

Note: this only implicates the existence of MC, not which variable causes MC.

Relationships between the two measures

Suppose that X is standardised to have mean 0 and variance 1, and we decompose $(X^\top X)^{-1}$ into $G \text{diag}(1/\lambda_1, \dots, 1/\lambda_p) G^\top$, then:

$$\begin{pmatrix} VIF_1 \\ \vdots \\ VIF_p \end{pmatrix} = \begin{pmatrix} g_{11}^2 & \cdots & g_{1p}^2 \\ \vdots & \ddots & \vdots \\ g_{p1}^2 & \cdots & g_{pp}^2 \end{pmatrix} \begin{pmatrix} \tau_1 \\ \vdots \\ \tau_p \end{pmatrix} = (G \circ G) \boldsymbol{\tau},$$

where $\tau_j = 1/\lambda_j$, $j = 1, \dots, p$.

Larger τ_p value indicates larger MC.

- It will be great if we have a formula of the form $\tau_p = f(VIF_1, \dots, VIF_p)$ to reveal the relationship between every variable x_j and the cause of MC, τ_p .

```
solve(eigen$u * eigen$u)[1:2,1:5]
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] -3.549883e+14  3.346813e+14  1.009966e+15 -6.696969e+13  1.492518e+14
## [2,] -1.050852e+13  9.907388e+12  2.989748e+13 -1.982468e+12  4.418220e+12
```

The mcvis method

We perform linear regression between τ_p and every VIF.

- By quantifying the linearity between τ_p and VIFs, we can diagnose MC-causing variables.
- How can we generate multiple "observations" of both τ_p and VIFs?
- Sampling!

VIF_1, \dots, VIF_p τ_1, \dots, τ_p

Bootstrap 1

$$VIF_1, \dots, VIF_p$$

$$\tau_1, \dots, \tau_p$$

Bootstrap 100

$$VIF_1, \dots, VIF_p$$

$$\tau_1, \dots, \tau_p$$

Bootstrap 200

$$VIF_1, \dots, VIF_p$$

$$\tau_1, \dots, \tau_p$$

.

.

.

.

.

.

Bootstrap 1000

$$VIF_1, \dots, VIF_p$$

$$\tau_1, \dots, \tau_p$$

Perform linear regression
extract t-statistic

Bootstrap 1

VIF_1, \dots, VIF_p

τ_1, \dots, τ_p

Bootstrap 100

VIF_1, \dots, VIF_p

τ_1, \dots, τ_p

Bootstrap 200

VIF_1, \dots, VIF_p

τ_1, \dots, τ_p

•
•
•

•
•
•

Bootstrap 1000

VIF_1, \dots, VIF_p

τ_1, \dots, τ_p

Perform linear regression
extract t-statistic

Bootstrap 1

VIF_1, \dots, VIF_p

τ_1, \dots, τ_p

Bootstrap 100

VIF_1, \dots, VIF_p

τ_1, \dots, τ_p

Bootstrap 200

VIF_1, \dots, VIF_p

τ_1, \dots, τ_p

•
•
•


•
•
•

Bootstrap 1000

VIF_1, \dots, VIF_p

τ_1, \dots, τ_p

Perform linear regression
extract t-statistic



Bootstrap 1

}

$K = 1$

$t_{1,1}, \dots, t_{p,1}$

Bootstrap 100

}

$K = 2$

$t_{1,2}, \dots, t_{p,2}$

Bootstrap 200


}

$K = 10$

$t_{1,10}, \dots, t_{p,10}$

Bootstrap 1000

$$\overline{t_j^2} = \left(\sum_{k=1}^K t_{j,k}^2 \right) / K$$



$$\overline{t_j^2} = \left(\sum_{k=1}^K t_{j,k}^2 \right) / K$$

$$\overline{t_1^2}, \overline{t_2^2}, \dots, \overline{t_p^2}$$

$$MC_j = \frac{\overline{t_j^2}}{\sum_{j=1}^p \overline{t_j^2}}$$

The `mcvis` package

1. MC-index


```
library(mcvis)
set.seed(13)
p = ncol(X)
mcvis_result = mcvis(X[, -p])
round(mcvis_result$MC[p-1, ], 2)
```

```
## log_runs log_outs log_ave log_fours log_sixes log_ducks log_hs
##      0.69      0.16      0.14      0.00      0.00      0.00      0.00
```

2. MC visualisation

```
ggplot_mcvis(mcvis_result)
```

Multi-collinearity plot



3. Shiny app for interactive exploration of data


Show entries

Search:

	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
1	1	50	0.096	2	0.1	0.15	1.9	-5.8	4.1	9.9	-0.31	0.16	0.29
2	2	50	-0.28	2.7	-0.35	-0.37	3.3	-5.4	6.9	12	0.34	-0.29	0.39
3	3	50	-0.09	2.4	-0.15	-0.082	2.1	-5.7	8.5	14	0.43	1.7	0.35
4	4	50	-0.51	3.2	-0.85	-0.44	2.5	-10	5.5	16	-0.31	0.47	0.45
5	5	50	-0.43	3.1	-0.45	-0.39	3	-9.3	6.6	16	-0.2	0.54	0.44
6	6	50	-0.4	3.9	-0.6	-0.36	4.1	-9.3	8.1	17	-0.094	-0.68	0.55
7	7	50	-0.2	2.6	0.15	-0.082	2.7	-6.7	5.5	12	-0.39	-0.15	0.36
8	8	50	1	5.6	1.1	0.8	6.3	-9.4	18	27	0.49	0.2	0.79
9	9	50	0.53	2.9	0.5	0.49	3.6	-4.8	5.8	11	0.12	-0.97	0.4

Showing 1 to 9 of 9 entries


Previous Next



Extension work: Multiple τ 's

```
ggplot_mcvis(mcvis_result, eig_max = 7)
```

Multi-collinearity plot








Simulated data from `mpLot`

- The R package `mpLot` (Tarr et. al. 2018) introduces a simulated data with correlated variables in context of variable selection.
- `mcvis` clearly identifies the correct cause of linearity (x_8) whereas VIF identifies (x_6).

	x_1	x_2	x_3	x_4	x_5	x_6	x_7	x_8	x_9
MC index	0.00885	0.0113	0.282	0.002	0.023	0.276	0.015	0.362	0.020
VIF	23.84	23.36	109.76	7.82	41.52	167.07	31.37	145.89	41.12

Final remarks

- mcvis provides a new MC-index and a visualisation of multicollinearity in linear regression.
- mcvis builds on top of classical statistics under a resampling framework and uncovers new sources of collinearity with an understanding of variability.
- Learn more from:
 -  [leaffur/mcvis](#)
 -  [kevinwang09/mcvispy](#)
 -  samuel.mueller@sydney.edu.au
 -  [@KevinWang009](#) and [@SamuelMuller74](#)



Bibliography

Marquardt, D. W. (1970). "Generalized Inverses, Ridge Regression, Biased Linear Estimation, and Nonlinear Estimation". In: Technometrics 12.3, pp. 591–612.

Belsley, D. A. (1984). "Demeaning Conditioning Diagnostics through Centering". In: The American Statistician 38.2, pp. 73–77.

Stewart, G. W. (1987). "Collinearity and Least Squares Regression". In: Statistical Science 2.1, pp. 68–84.

O'Brien, R. M. (2007). "A Caution Regarding Rules of Thumb for Variance Inflation Factors". In: Quality & Quantity 41.5, pp. 673–690.

Sheather, S. (2009). A Modern Approach to Regression with R. Springer Texts in Statistics. New York, NY: Springer New York.

Friendly, M. and E. Kwan (2009). "Statistical computing and graphics where's Waldo? Visualizing collinearity diagnostics". In: The American Statistician 63.1, pp. 56–65.

list()